

MusicGAN

A GAN capable of composing
little pieces of music

Dataset



Type of data


- Great collection of piano rolls
- ~20000 different Multitrack objects
- For each beat of the song is known what notes are being played
- Five different instruments


Preprocessing



Reduced note range

- Most notes in a few octaves
- Improvement in memory and time
- Little loss of information






Normalization

- First experiments **sounded weird**
- Analysis showed different **note distribution** from dataset
- Train only on a **single scale**
- **Music theory** knowledge needed
- Source of claims is a **domain expert**




Major scales

- Typical note **pattern**
- All equal, but translated by a **certain offset**
- They sound the same to **99.99%** of people
- Used in **75%** of dataset



Major or natural minor?

- Same notes but with different roles
- Possible to distinguish with clustering
- Natural minor not really used in practice



Other scales


- Melodic/harmonic minor, pentatonic, hexatonic, blues, the list goes on
- Little representation for single type
- Decided to train on major scales

What was done

- Look at **note frequency** and determine scale
- If **major scale**, transpose every note, so that it becomes a C major song
- If **other type** of scale, discard

Approach






G

e

C


C



Wasserstein Metric

$$W(\mathbb{P}_r, \mathbb{P}_g) = \min_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma}[\|x - y\|]$$

Wasserstein Metric



Reference: Wasserstein GAN, Arjovsky et al.

Algorithm

Require: α , the learning rate. m , the batch size. n_{critic} , the number of iterations of the critic per generator iteration. β , the step-size of the gradient penalty.

Inputs: w_0 , initial critic parameters. θ_0 , initial generator parameters.

```
1: while  $\theta$  has not converged do
2:   for  $t=0, \dots, n_{critic}$  do
3:     Sample  $\{x^{(i)}\}_{i=1}^m \sim \mathbb{P}$  a batch from the real data
4:     Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples
5:     
$$g_w \leftarrow \nabla_w \left[ \frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right]$$

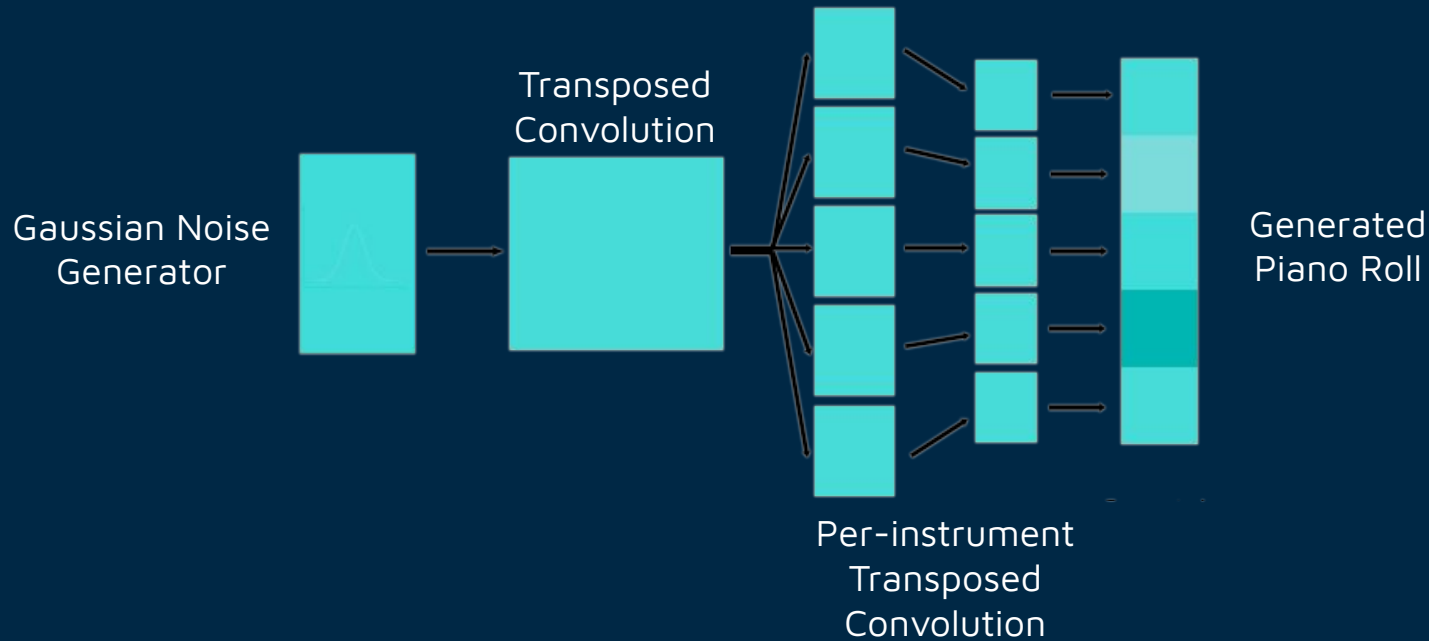
6:     
$$w \leftarrow w + \alpha \cdot \text{RMSProp}(w, g_w) - \beta \cdot \nabla_w g_w$$

7:   end for
8:   Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples
9:   
$$g_\theta \leftarrow \nabla_\theta \left[ \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right]$$

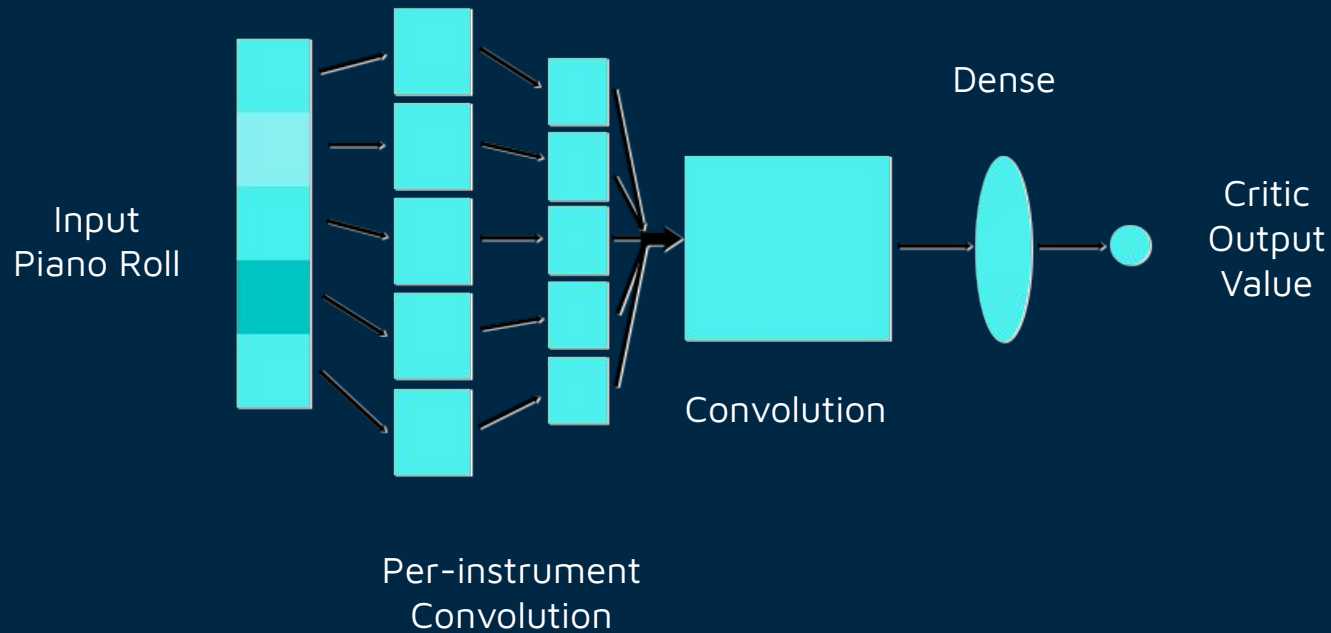
10:  
$$\theta \leftarrow \theta + \alpha \cdot \text{RMSProp}(\theta, g_\theta)$$

11: end while
```

Generator Architecture



Critic Architecture



Training & Results



What we expect:


- Not an hit
- A catchy musical motif
- No dissonances and no random notes

Training details:

- “Supervised” training
- Google Colab for memory needs
- Random noise taken from a Gaussian Distribution
- 1:5 ratio of weights update between G and C

First attempt:

Give all the dataset to the network




The critic is **confused**, there is no clarity in the dataset: its value is **meaningless**


Second attempt:

Helping the networks by training on **one genre**

Increased data **consistency**



The critic is more **reliable**




The generator work is facilitated
(WGAN)

+

Diversification of the generated
songs



Results analysis:



At the minimum of the G loss:

- Notes are **less piled up**
- **Patterns** became more definite
- Instruments are **synchronized**
- **Consistent** pauses